

RUN vs CMD vs ENTRYPOINT

Some Docker instructions look similar and cause confusion among developers who just started using Docker or do it irregularly. In this post I will explain the difference between CMD, RUN, and ENTRYPOINT on examples.

- **RUN** executes command(s) in a new layer and creates a new image. E.g., it is often used for installing software packages.
- **CMD** sets default command and/or parameters, which can be overwritten from command line when docker container runs.
- **ENTRYPOINT** configures a container that will run as an executable.

Docker images and layers

When Docker runs a container, it runs an image inside it. This image is usually built by executing Docker instructions, which add layers on top of existing image or OS distribution. OS distribution is the initial image and every added layer creates a new image.

Final Docker image reminds an with OS distribution inside and a number of layers on top of it. For example, your image can be built by installing a number of deb packages and your application on top of Ubuntu 14.04 distribution.

Shell and Exec forms

All three instructions (RUN, CMD and ENTRYPOINT) can be specified in **shell** form or **exec** form. Let's get familiar with these forms first, because the forms usually cause more confusion than instructions themselves.

Shell form

<instruction> <command>

Examples:Dockerfile

FROM ubuntu

RUN apt-get update -y

RUN apt-get install python3 -y

CMD echo "Hello world"

ENTRYPOINT echo "Hello world"

build the docker image \$ docker build -t test .

When instruction is executed in shell form it calls `/bin/sh -c <command>` under the hood and normal shell processing happens. For example, the following snippet in Dockerfile

Ex: Dockerfile

FROM ubuntu

RUN apt-get update -y

ENV name madhu sudhan

ENTRYPOINT echo "Hello, \$name"

\$ docker build -t test1 .

when container runs as **docker run -it test1** will produce output

```
[ec2-user@ip-172-31-34-155 ~]$ docker run -it test1
Hello, madhu sudhan
```

Note that variable name is replaced with its value.

Exec form

This is the preferred form for CMD and ENTRYPOINT instructions.

<instruction> ["executable", "param1", "param2", ...]

Examples:

RUN ["apt-get", "install", "python3"]

CMD ["/bin/echo", "Hello world"]

ENTRYPOINT ["/bin/echo", "Hello world"]

When instruction is executed in exec form it calls executable directly, and shell processing does not happen. For example, the following snippet in **Dockerfile**

FROM ubuntu

RUN apt-get update -y

ENV name madhu sudhan

ENTRYPOINT ["/bin/echo", "Hello, \$name"]

\$ docker build -t test2 .

when container runs as **docker run -it test2** will produce output

```
[ec2-user@ip-172-31-34-155 ~]$ docker run -it test2
Hello, $name
[ec2-user@ip-172-31-34-155 ~]$
```

Note that variable name is not substituted.

How to run bash?

If you need to run bash (or any other interpreter but sh), use exec form with /bin/bash as executable. In this case, normal shell processing will take place. For example, the following snippet in **Dockerfile**

FROM ubuntu

RUN apt-get update -y

ENV name madhu sudhan

ENTRYPOINT ["/bin/bash", "-c", "echo Hello, \$name"]

\$ **docker build -t test3 .**

when container runs as **docker run -it test3** will produce output

```
[ec2-user@ip-172-31-34-155 ~]$ docker run -it test3
Hello, madhu sudhan
[ec2-user@ip-172-31-34-155 ~]$
```

RUN

RUN instruction allows you to install your application and packages required for it. It executes any commands on top of the current image and creates a new layer by committing the results. Often you will find multiple RUN instructions in a Dockerfile.

RUN has **two** forms:

RUN <command> (shell form)

RUN ["executable", "param1", "param2"] (exec form)

(The forms are described in detail in Shell and Exec forms section above.)

A good illustration of RUN instruction would be to install multiple version control systems packages:

**RUN apt-get update && apt-get install -y **

**bzr **

**cvs **

**git **

**mercurial **

subversion

Note that apt-get update and apt-get install are executed in a single RUN instruction. This is done to make sure that the latest packages will be installed. If apt-get install were in a separate RUN instruction, then it would reuse a layer added by apt-get update, which could have been created a long time ago.

CMD

CMD instruction allows you to set a default command, which will be executed only when you run container without specifying a command. If Docker container runs with a command, the default command will be ignored. If Dockerfile has more than one CMD instruction, all but last CMD instructions are ignored.

CMD has three forms:

CMD ["executable", "param1", "param2"] (exec form, preferred)

CMD ["param1", "param2"] (sets additional default parameters for ENTRYPOINT in exec form)

CMD command param1 param2 (shell form)

Again, the first and third forms were explained in Shell and Exec forms section. The second one is used together with ENTRYPOINT instruction in exec form. It sets default parameters that will be added after ENTRYPOINT parameters if container runs without command line arguments. See ENTRYPOINT for example.

Let's have a look how CMD instruction works. The following snippet in **Dockerfile**

```
FROM ubuntu
```

```
RUN apt-get update -y
```

```
CMD echo "Hello world"
```

```
$ docker build -t test4 .
```

when container runs as **docker run -it test4** will produce output

```
[ec2-user@ip-172-31-34-155 ~]$ docker run -it test4
Hello world
[ec2-user@ip-172-31-34-155 ~]$
```

but when container runs with a command, e.g., **docker run -it <image> /bin/bash**, CMD is ignored and bash interpreter runs instead:

```
$ docker run -it test4 /bin/bash
```

```
root@7de4bed89922:/#
```

ENTRYPOINT

ENTRYPOINT instruction allows you to configure a container that will run as an executable. It looks similar to CMD, because it also allows you to specify a command with parameters. The difference is ENTRYPOINT command and parameters are not ignored when Docker container runs with command line parameters. (There is a way to ignore ENTRYPOINT, but it is unlikely that you will do it.)

ENTRYPOINT has **two** forms:

ENTRYPOINT ["executable", "param1", "param2"] (exec form, preferred)

ENTRYPOINT command param1 param2 (shell form)

Be very careful when choosing ENTRYPOINT form, because forms behaviour differs significantly.

Exec form

Exec form of ENTRYPOINT allows you to set commands and parameters and then use either form of CMD to set additional parameters that are more likely to be changed. ENTRYPOINT arguments are always used, while CMD ones can be overwritten by command line arguments provided when Docker container runs. For example, the following snippet in **Dockerfile**

```
FROM ubuntu
```

```
RUN apt-get update -y
```

```
ENTRYPOINT ["/bin/echo", "Hello"]
```

```
CMD ["world"]
```

```
$ docker build -t test5 .
```

when container runs as **docker run -it test5** will produce output

```
[ec2-user@ip-172-31-34-155 ~]$ docker run -it test5
Hello world
[ec2-user@ip-172-31-34-155 ~]$
```

but when container runs as **docker run -it test5 Reddy** will result in

```
hello world
[ec2-user@ip-172-31-34-155 ~]$ docker run -it test5 Reddy
Hello Reddy
[ec2-user@ip-172-31-34-155 ~]$
```

Shell form

Shell form of ENTRYPOINT ignores any CMD or docker run command line arguments.

The bottom line

Use RUN instructions to build your image by adding layers on top of initial image.

Prefer ENTRYPOINT to CMD when building executable Docker image and you need a command always to be executed. Additionally use CMD if you need to provide extra default arguments that could be overwritten from command line when docker container runs.

Choose CMD if you need to provide a default command and/or arguments that can be overwritten from command line when docker container runs.